

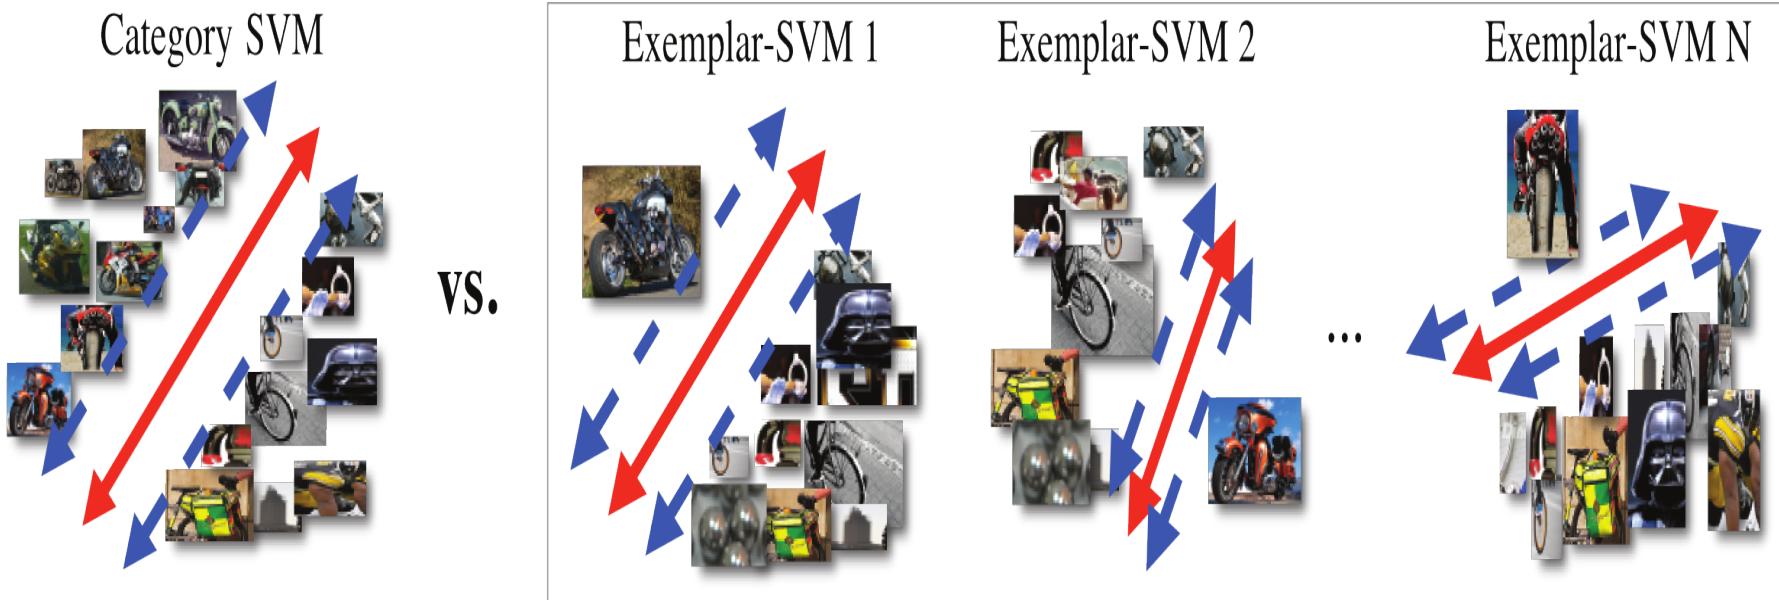
Ensemble of Exemplar-SVMs for Object Detection and Beyond

**T. Malisiewicz, A. Gupta, A. Efros,
ICCV 2011.**

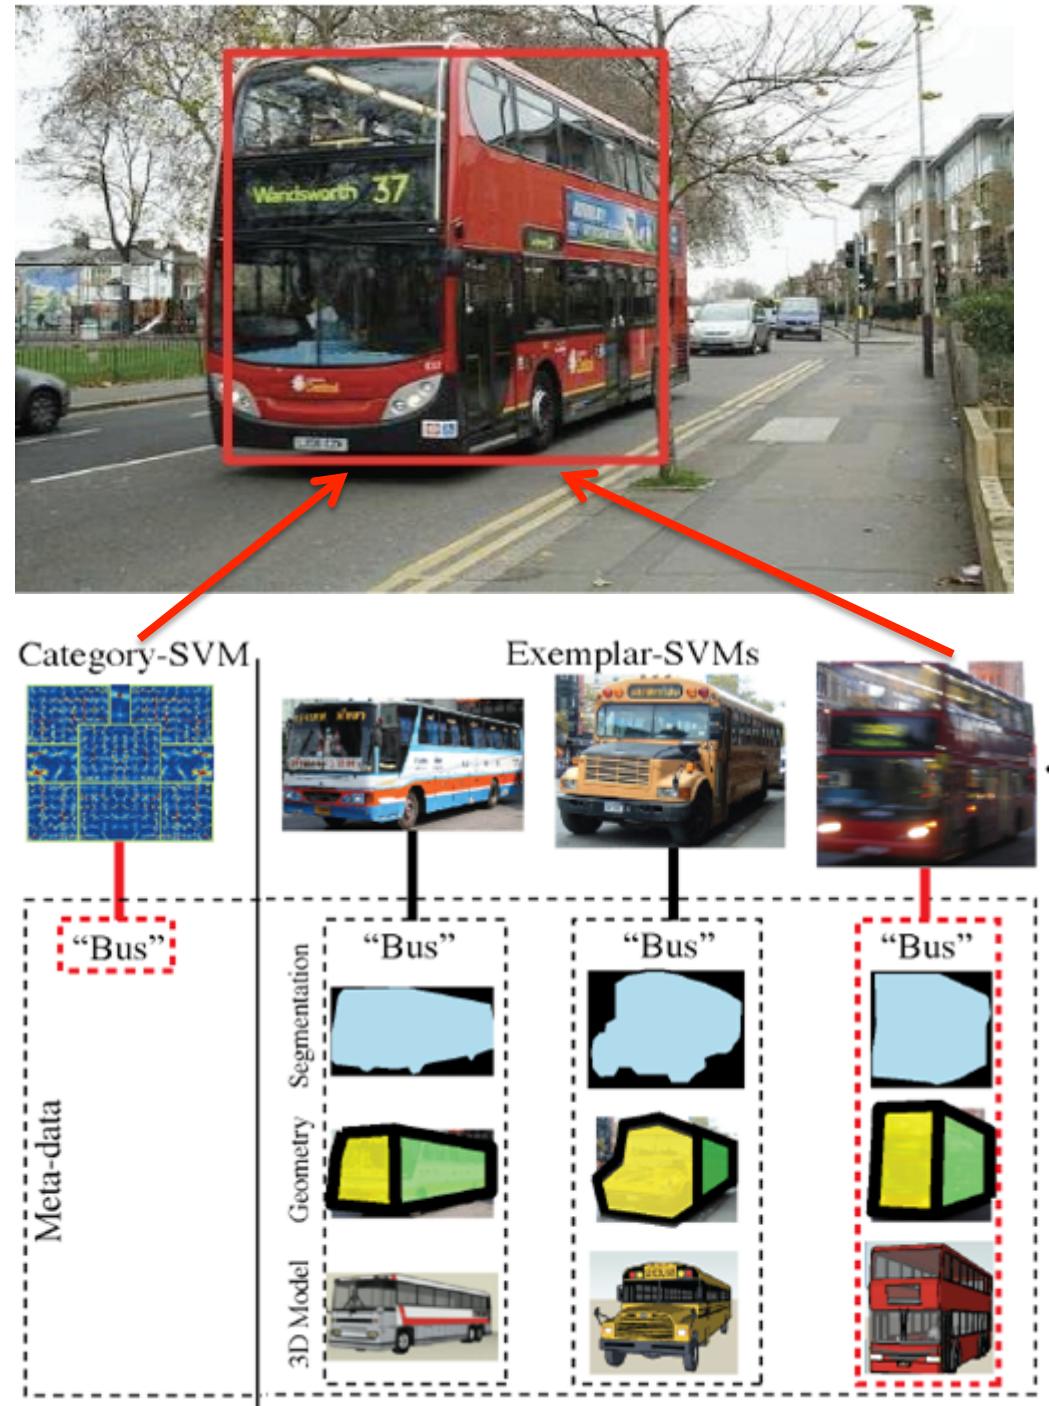
Abstract

- Motivation & Related Work
- Exemplar-SVMs
- Object Detection Results
- Beyond Detection

Ensemble of Exemplar-SVMs



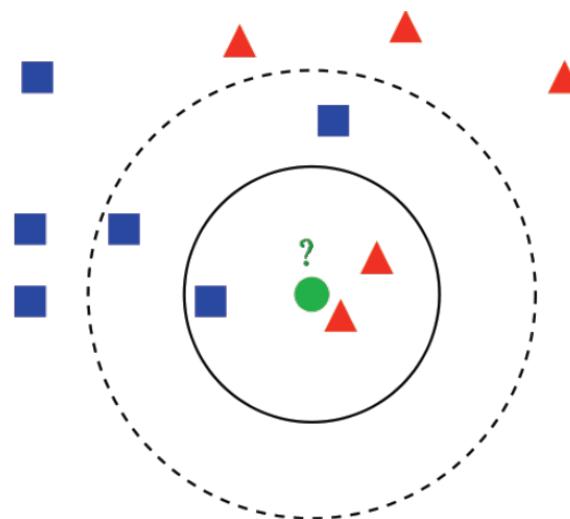
Motivation



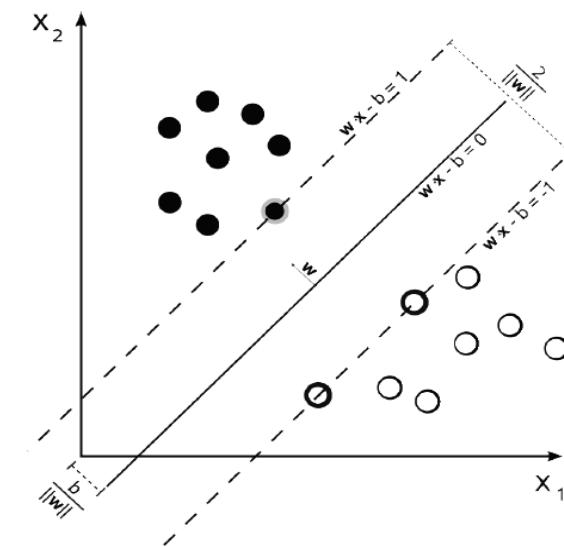
Motivation

- Idea of associating a new instance with something seen in the past
 - Exemplar theory in cognitive psychology
 - Case based reasoning in AI
 - Instance based methods in ML

Motivation - Exemplar Reasoning: Non-parametric



KNN: non-parametric



SVM: parametric

Exemplar Theory in CV

- Object Alignment
- Scene Recognition
- Image Parsing
- Object Detection(not competitive)
 - Why?

Parametric Approaches

- Can handle negative data well
- No need to explicitly store negative data(vs KNN)
- Positive Data?
 - Implicit assumption that all positive examples are visually related, which
 - Results in over generalized models

Not that similar!



Motivation-Desirable Approach

- Strengths of Dalal/Triggs/Felzenswalb/Ramanan style detector (Hog/DPM)
 - Powerful descriptor
 - Discriminative framework
 - Handle massive amount of negatives

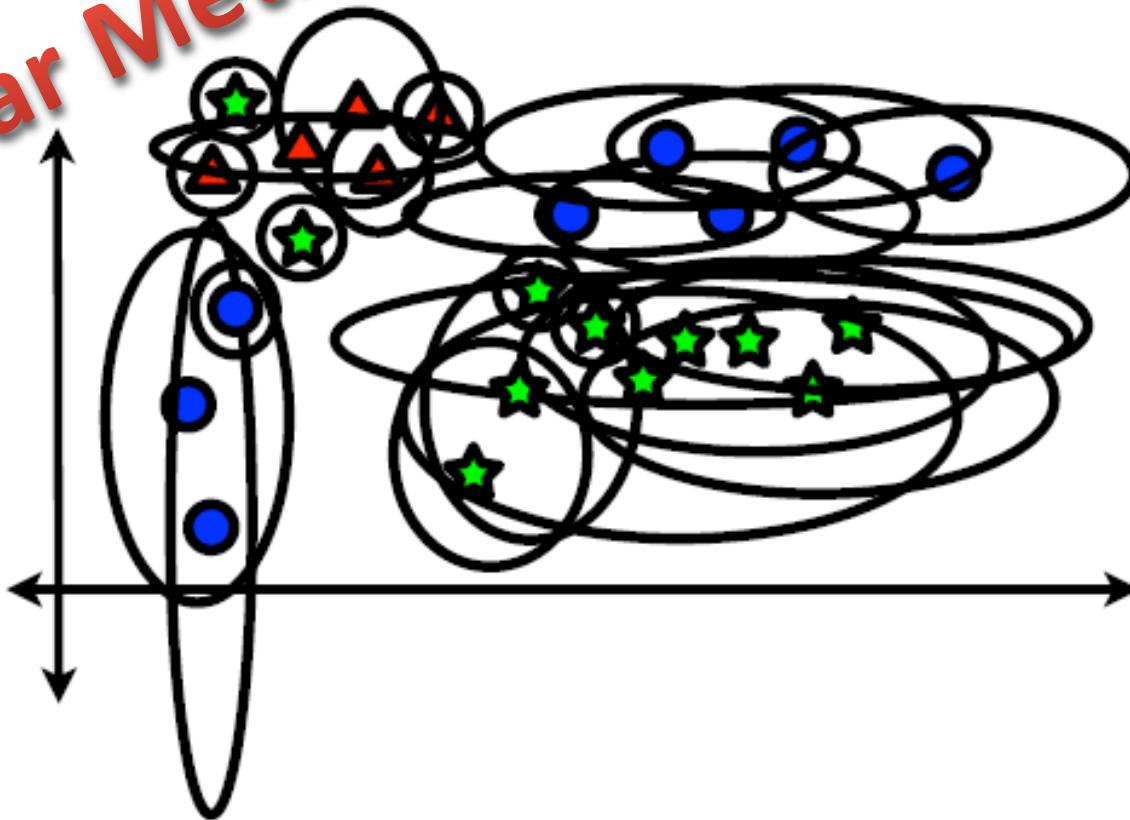
parametric
negatives

- Not rigidly representing positives

- Good association for meta-data transfer

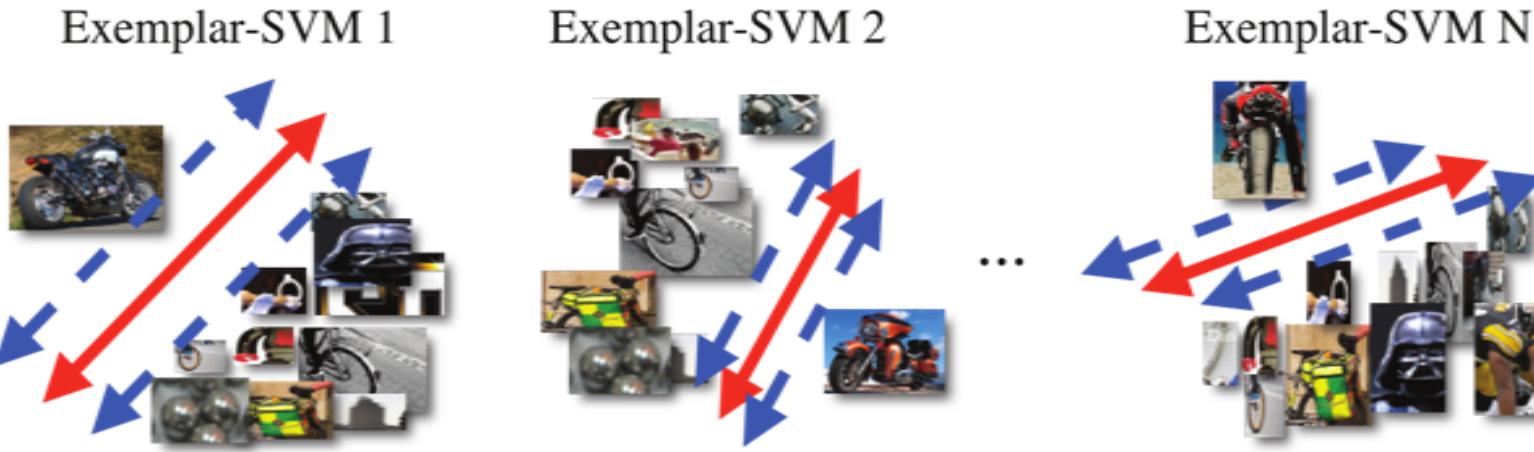
non-parametric positives

Per-Exemplar Methods



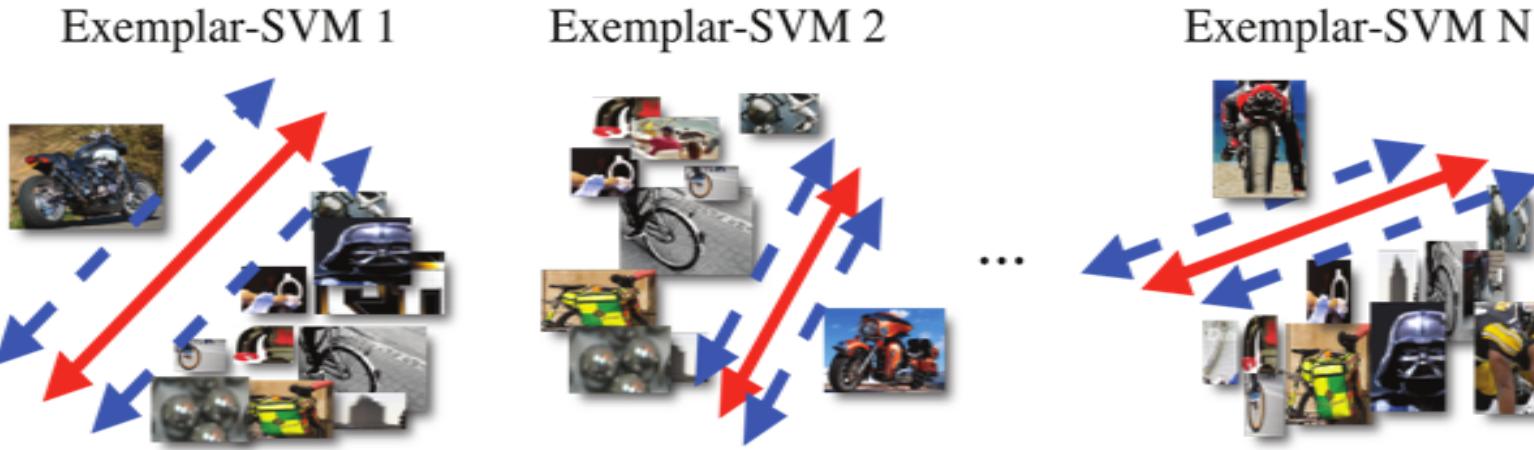
- NN-method: distance similarity function for each exemplar
- Better than a single similarity measure across all exemplars

Exemplar-SVMs



- Effectiveness of discriminatively trained object detectors
- Explicit correspondence of Nearest Neighbor approaches

Exemplar-SVMs



- Learn a model for each **positive** with Linear SVM
- Each **Exemplar-SVM** is trained with **single positive instance**
- **Exemplar's decision boundary is defined, by what it's not!**

Exemplar-SVMs

Exemplar-SVM 1



Exemplar-SVM 2



Exemplar-SVM N



- Because each exemplar is defined by a single positive instance, different features can be used for each exemplar
- Adapt features to each exemplar's aspect ratio



7x4 HOG



4x8 HOG

Exemplar-SVMs

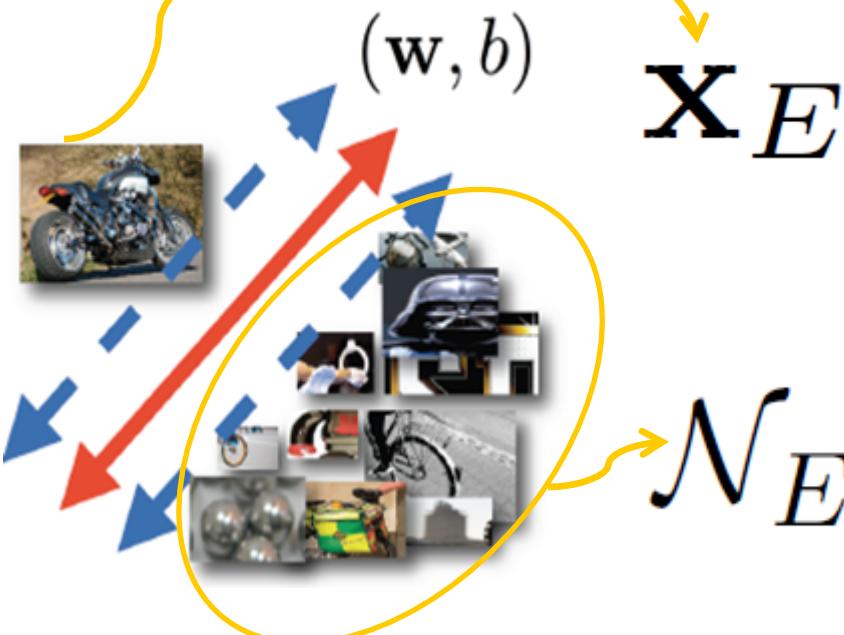
Exemplar E's Objective function:

$$\Omega_E(\mathbf{w}, b) = \|\mathbf{w}\|^2 + C_1 h(\mathbf{w}^T \mathbf{x}_E + b) + C_2 \sum_{\mathbf{x} \in \mathcal{N}_E} h(-\mathbf{w}^T \mathbf{x} - b)$$

$$C_1 = 0.5$$

$$C_2 = 0.1$$

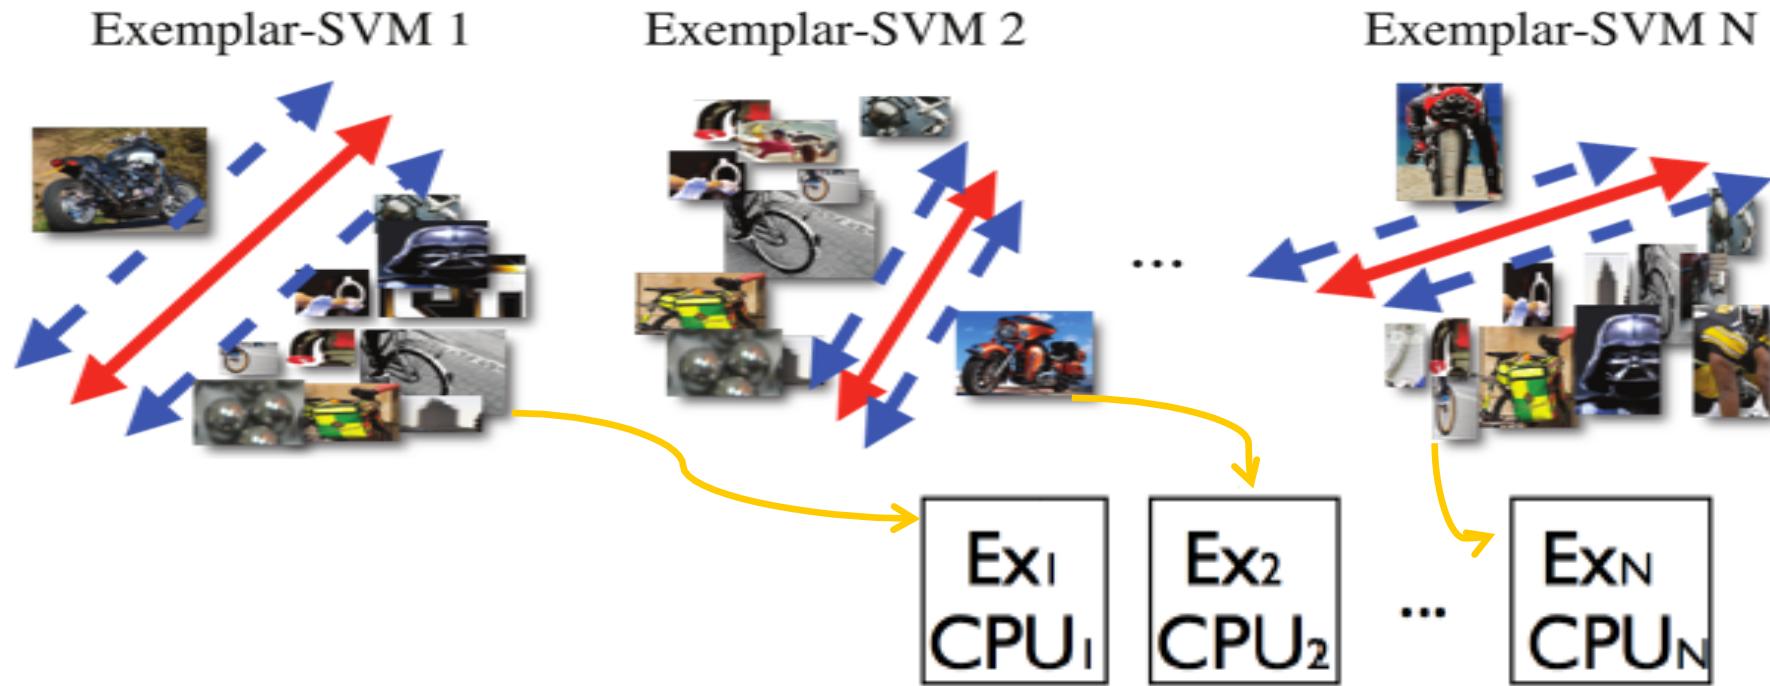
$$h(\mathbf{x}) = \max(1 - \mathbf{x}, 0) \text{ “hinge-loss”}$$



Exemplar represented by ~100 HOG cells

Windows from images not containing any in-class instances (~2,000 images x ~10,000 windows/image = ~2M negatives)

Large-Scale Training

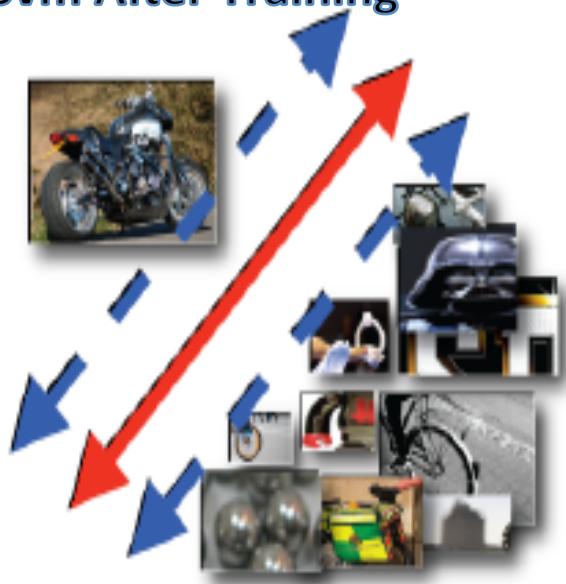


- Parallel training on clusters



Exemplar-SVM Calibration

Svm After Training



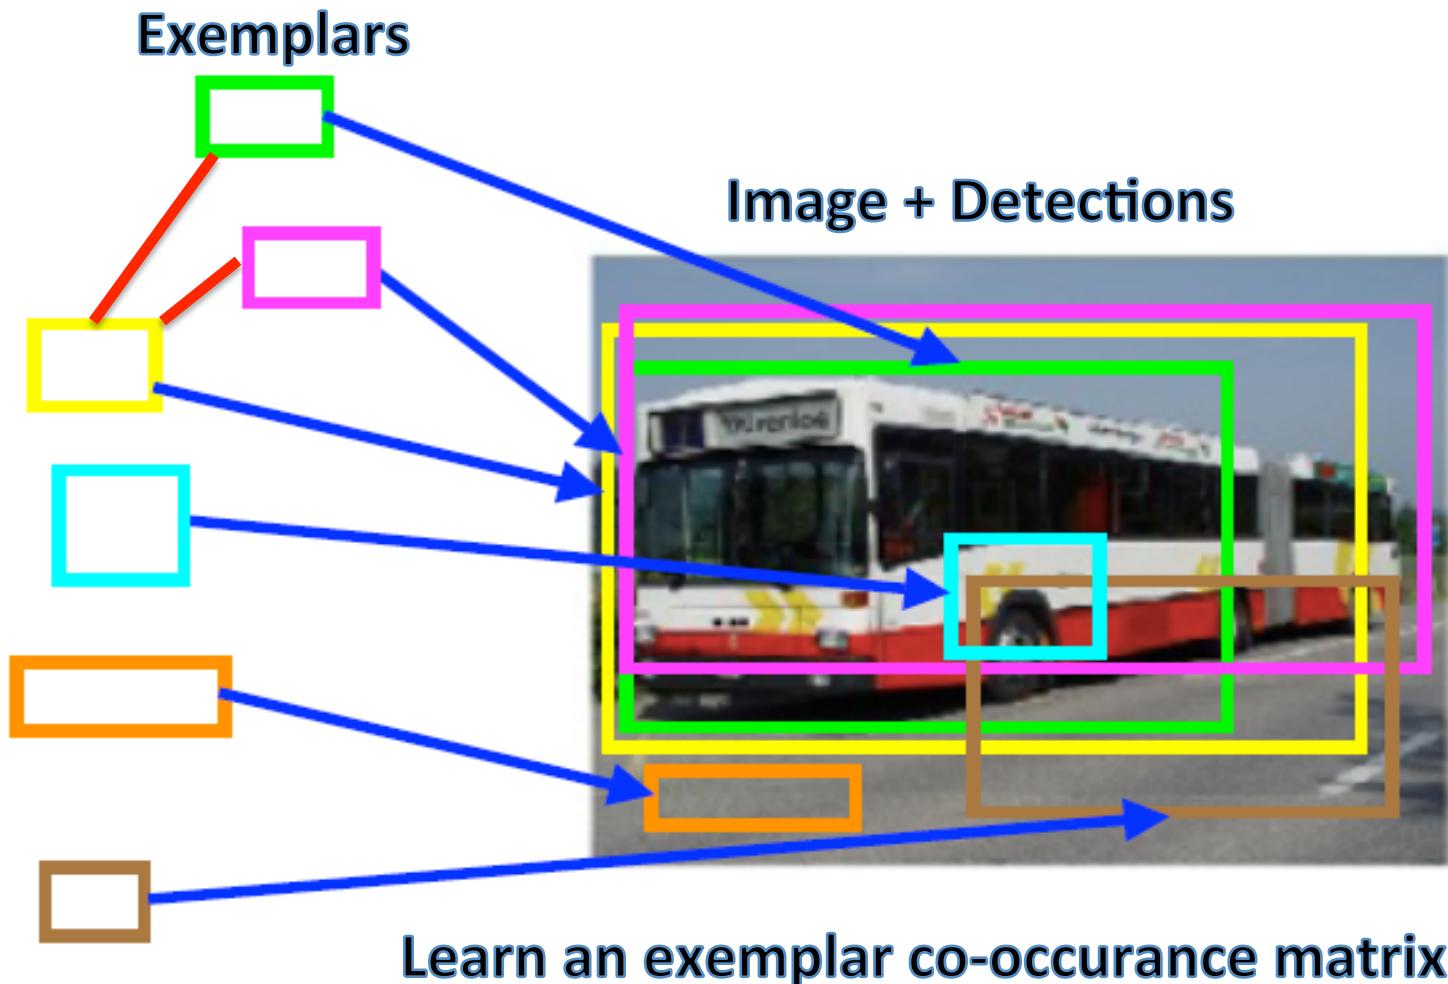
- Apply exemplar-SVM to held out negative images and all positive images
- Fit sigmoid to responses [Plat 1999]

Svm After Calibration

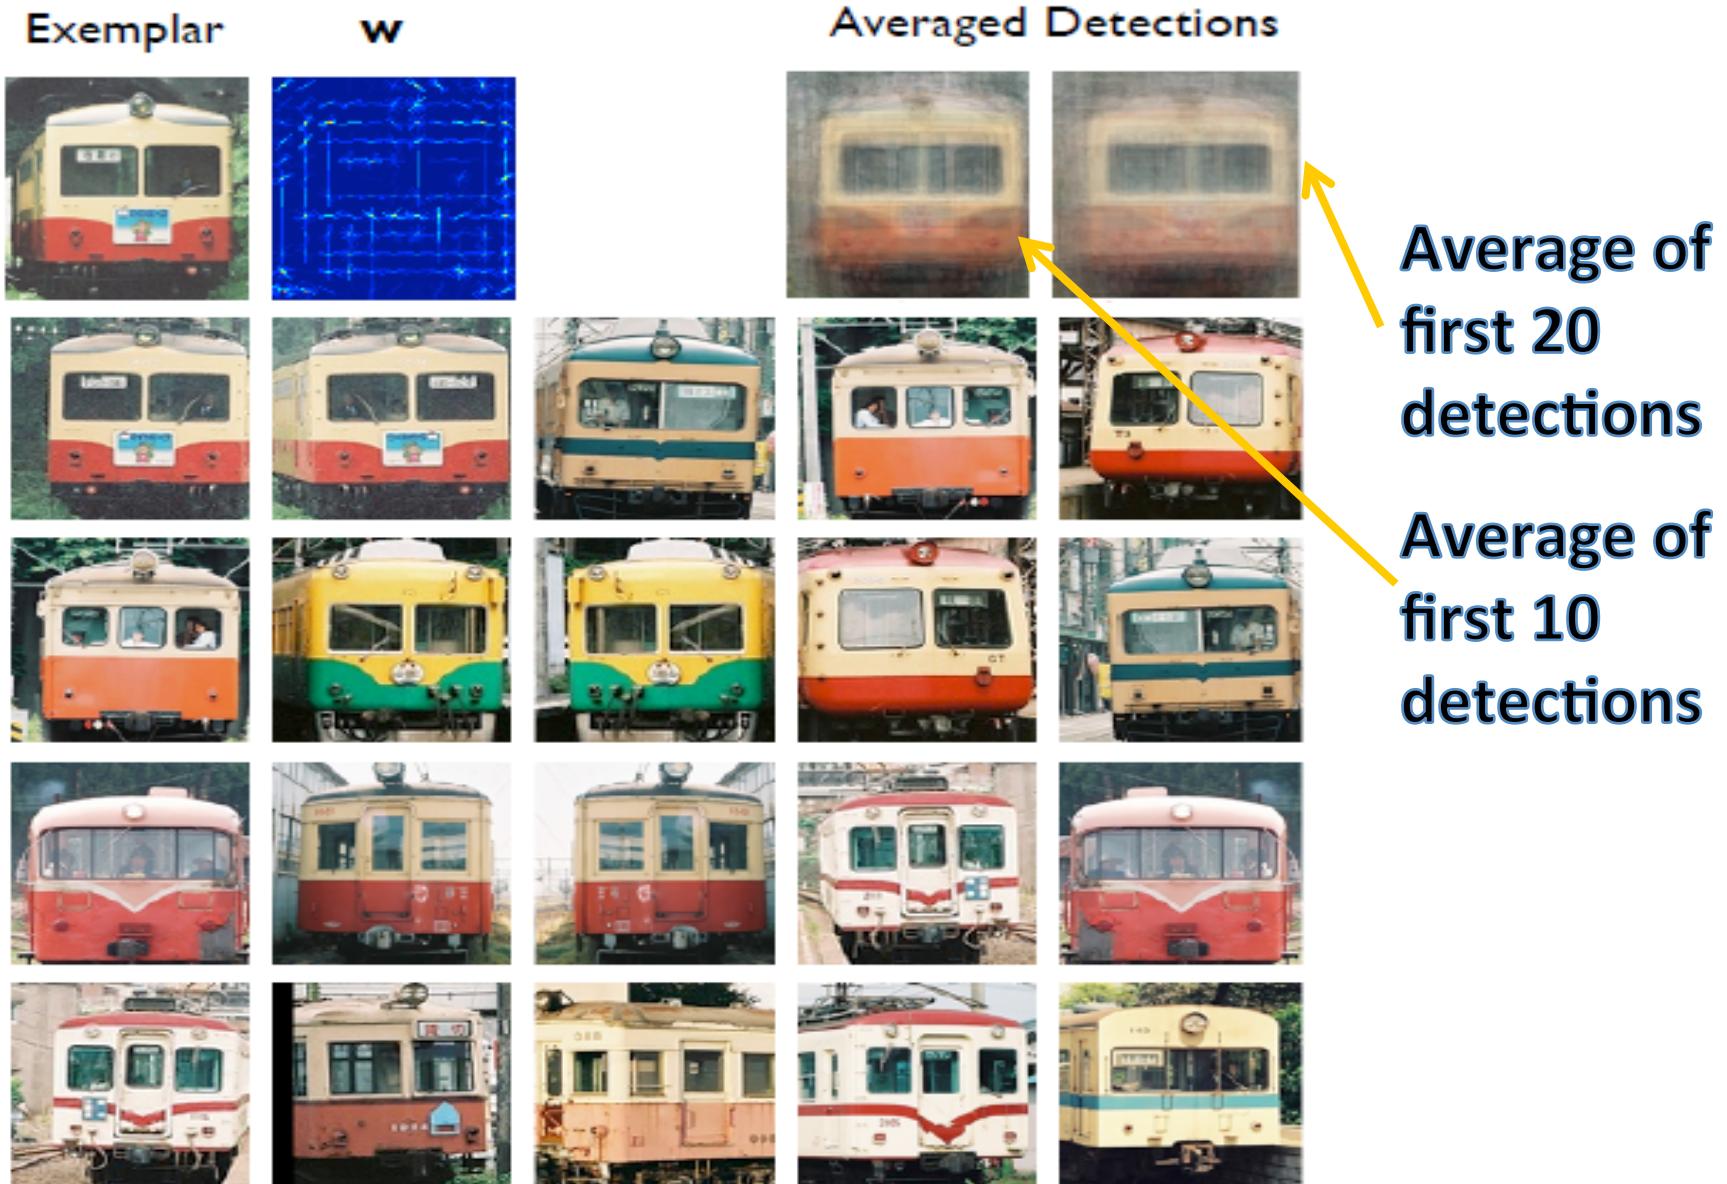


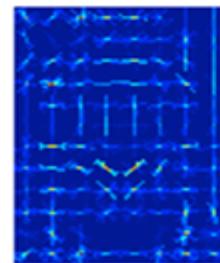
$$f(\mathbf{x}|\mathbf{w}_E, \alpha_E, \beta_E) = \frac{1}{1 + e^{-\alpha_E(\mathbf{w}_E^T \mathbf{x} - \beta_E)}}$$

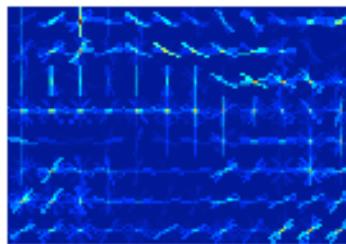
Ensemble of Exemplar-SVMs

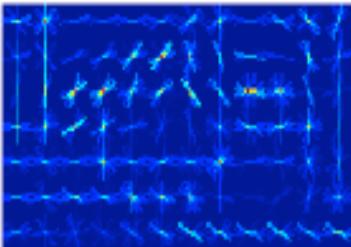


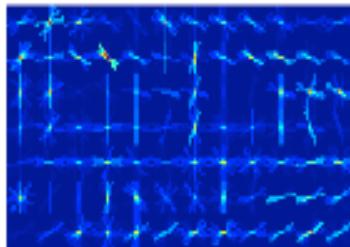
Qualitative Results











Evaluating Exemplar SVMs

- Nearest Neighbor
 - No Learning
- Per-Exemplar Distance Functions
 - Learning in distance-to-exemplar space
[Malisiewicz et. al. 2008]
- Exemplar SVMs

Comparison of 3 methods



*Learned Distance Function

Quantitative Results

- Pascal VOC 2007 dataset
- A standard computer vision object detection benchmark
- 20 object categories
- Machine performance is far below human

Pascal VOC: Object Category Detection Results

Approach	aeroplane	bicycle	bird	boat	bottle	bus	car	cat	chair	cow	diningtable	dog	horse	motorbike	person	pottedplant	sheep	sofa	train	tvmonitor	mAP
NN	.006	.094	.000	.005	.000	.006	.010	.092	.001	.092	.001	.004	.096	.094	.005	.018	.009	.008	.096	.144	.039
NN+Cal	.056	.293	.012	.034	.009	.207	.261	.017	.094	.111	.004	.033	.243	.188	.114	.020	.129	.003	.183	.195	.110
DFUN+Cal	.162	.364	.008	.096	.097	.316	.366	.092	.098	.107	.002	.093	.234	.223	.109	.037	.117	.016	.271	.293	.155
E-SVM+Cal	.204	.407	.093	.100	.103	.310	.401	.096	.104	.147	.023	.097	.384	.320	.192	.096	.167	.110	.291	.315	.198
E-SVM+Co-occ	.208	.480	.077	.143	.131	.397	.411	.052	.116	.186	.111	.031	.447	.394	.169	.112	.226	.170	.369	.300	.227
CZ [6]	.262	.409	-	-	-	.393	.432	-	-	-	-	-	.375	-	-	-	-	.334	-	-	
DT [7]	.127	.253	.005	.015	.107	.205	.230	.005	.021	.128	.014	.004	.122	.103	.101	.022	.056	.050	.120	.248	.097
LDPM [9]	.287	.510	.006	.145	.265	.397	.502	.163	.165	.166	.245	.050	.452	.383	.362	.090	.174	.228	.341	.384	.266

Table 1. PASCAL VOC 2007 object detection results. We compare our full system (ESVM+Co-occ) to four different exemplar based baselines including NN (Nearest Neighbor), NN+Cal (Nearest Neighbor with calibration), DFUN+Cal (learned distance function with calibration) and ESVM+Cal (Exemplar-SVM with calibration). We also compare our approach against global methods including our implementation of Dalal-Triggs (learning a single global template), LDPM [9] (Latent deformable part model), and Chum et al. [6]'s exemplar-based method. [The NN, NN+Cal and DFUN+Cal results for person category are obtained using 1250 exemplars]

Object Category Detection

mAP on PASCAL VOC 2007 detection task

NN + Cal	0.110
DFUN + Cal	0.155
Exemplar-SVMs + Cal	0.198
Exemplar-SVMs + Co-occ	0.227
DT*	0.097
LDPM**	0.266

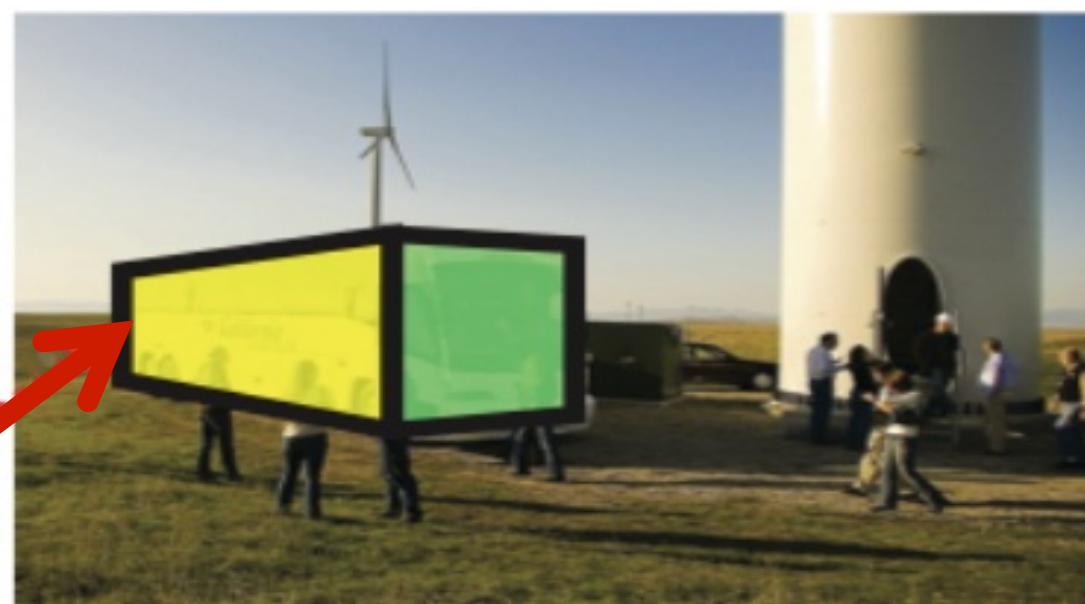
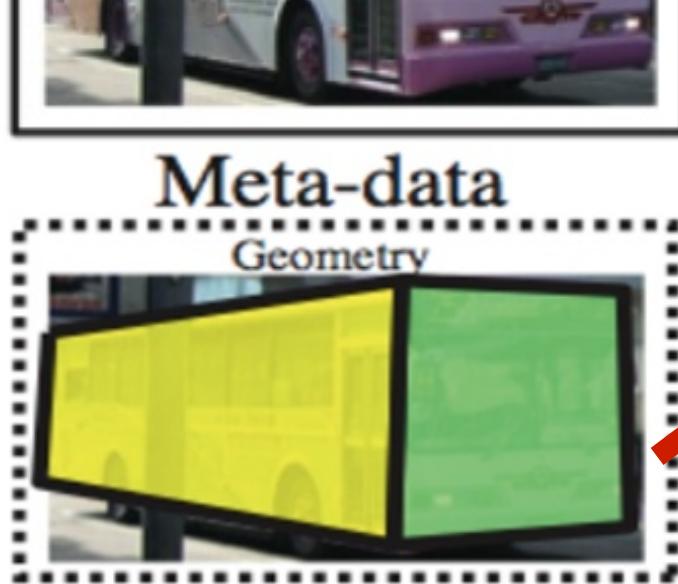
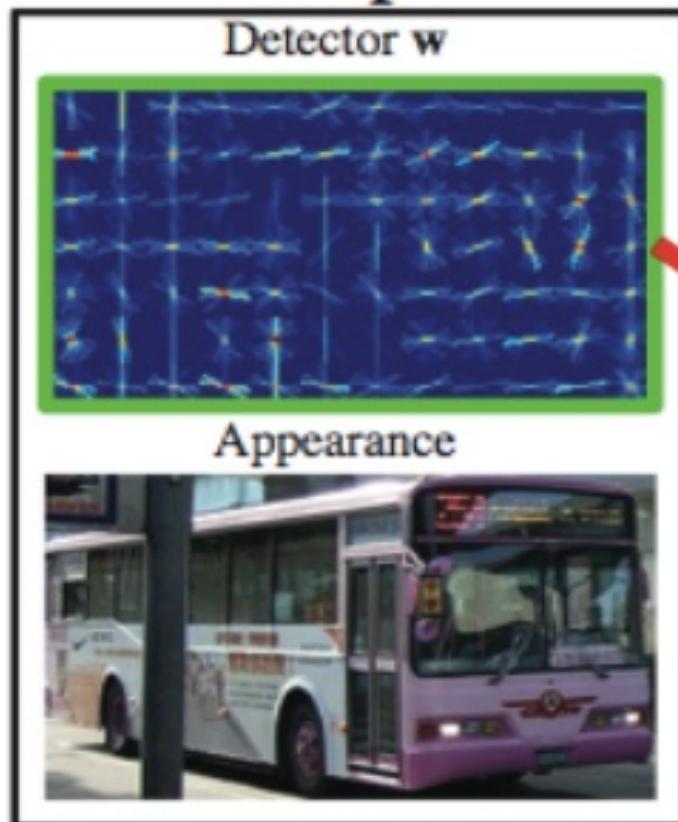
*Dalal et al. 2005

**Felzenszwalb et al. 2010

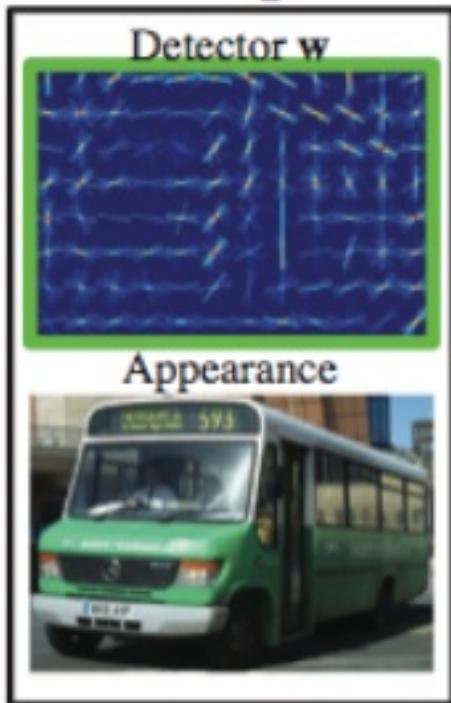
Beyond Object Detection

- Geometry Transfer
- 3D Model Transfer
- Relating Object Priming

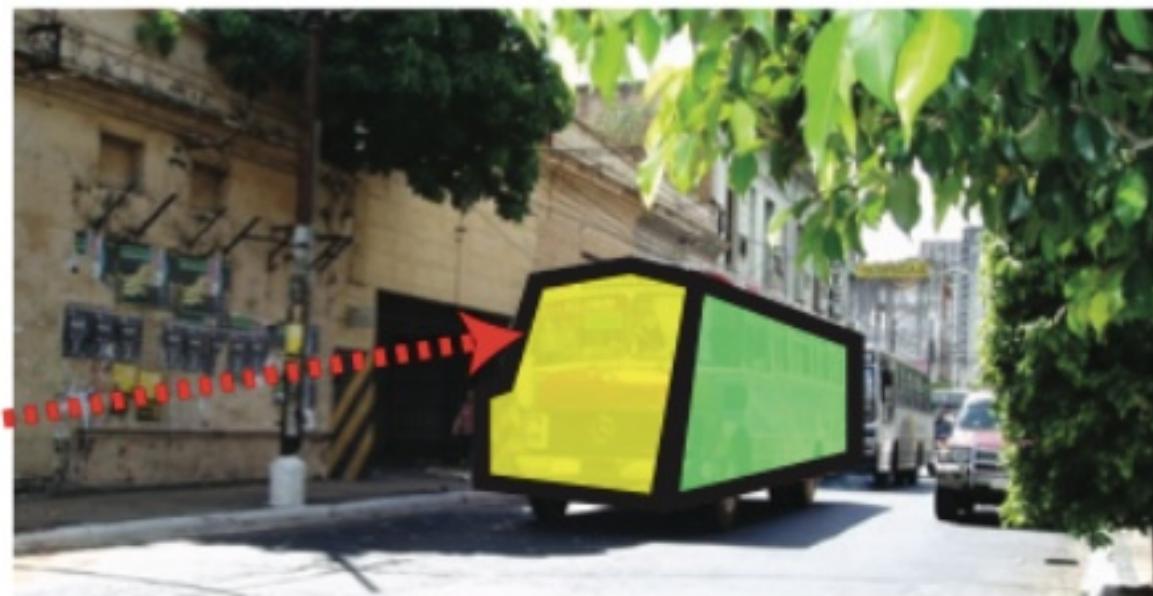
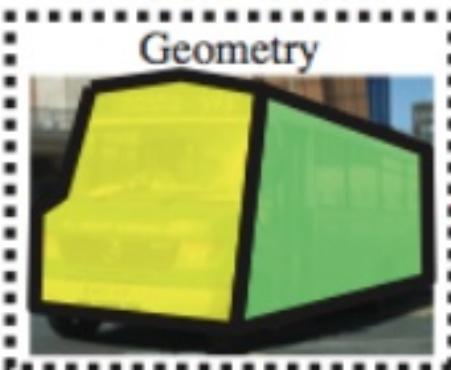
Exemplar



Exemplar



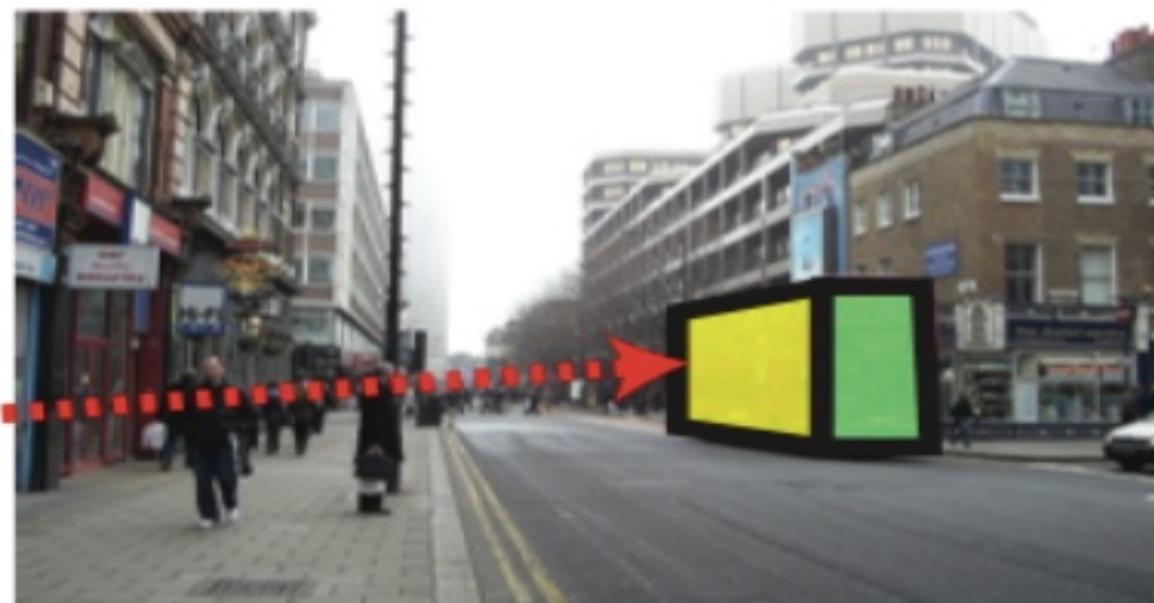
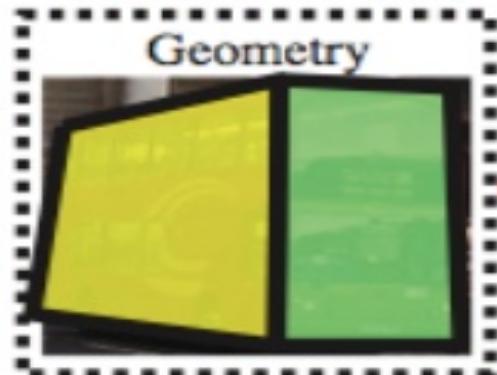
Meta-data



Exemplar



Meta-data



Geometry Transfer: Evaluation on Buses

- Measure pixelwise accuracy on the 3-class geometric-labeling problem:
“left”, “front”, “right” -facing
- 43.0% Hoeim et al. 2005
- 51.0% Category-SVM* + NN
- 62.3% Exemplar-SVMs

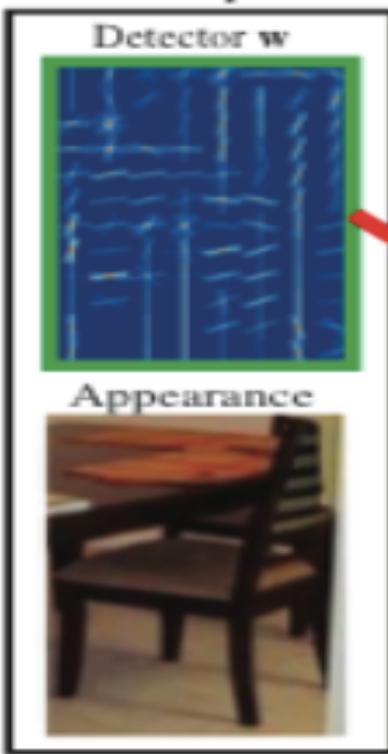
*felzenswalb et al. 2010

3D Model Transfer

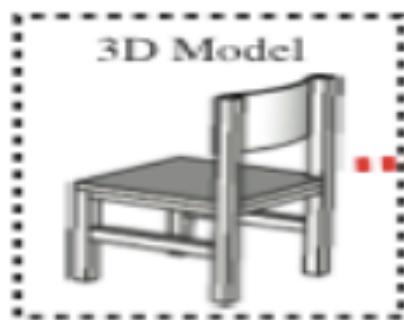


Manually align 3D
model from Google
3D Warehouse with
a subset of PASCAL
VOC “chair”
exemplars

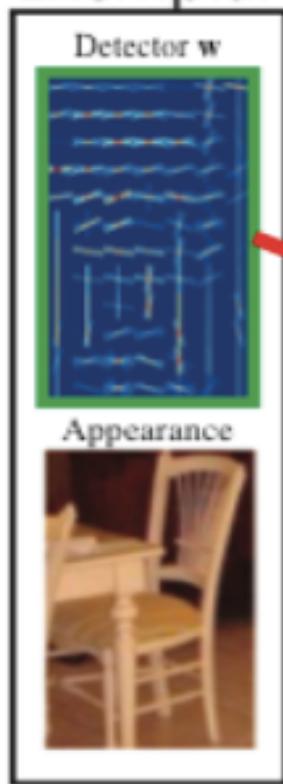
Exemplar



Meta-data



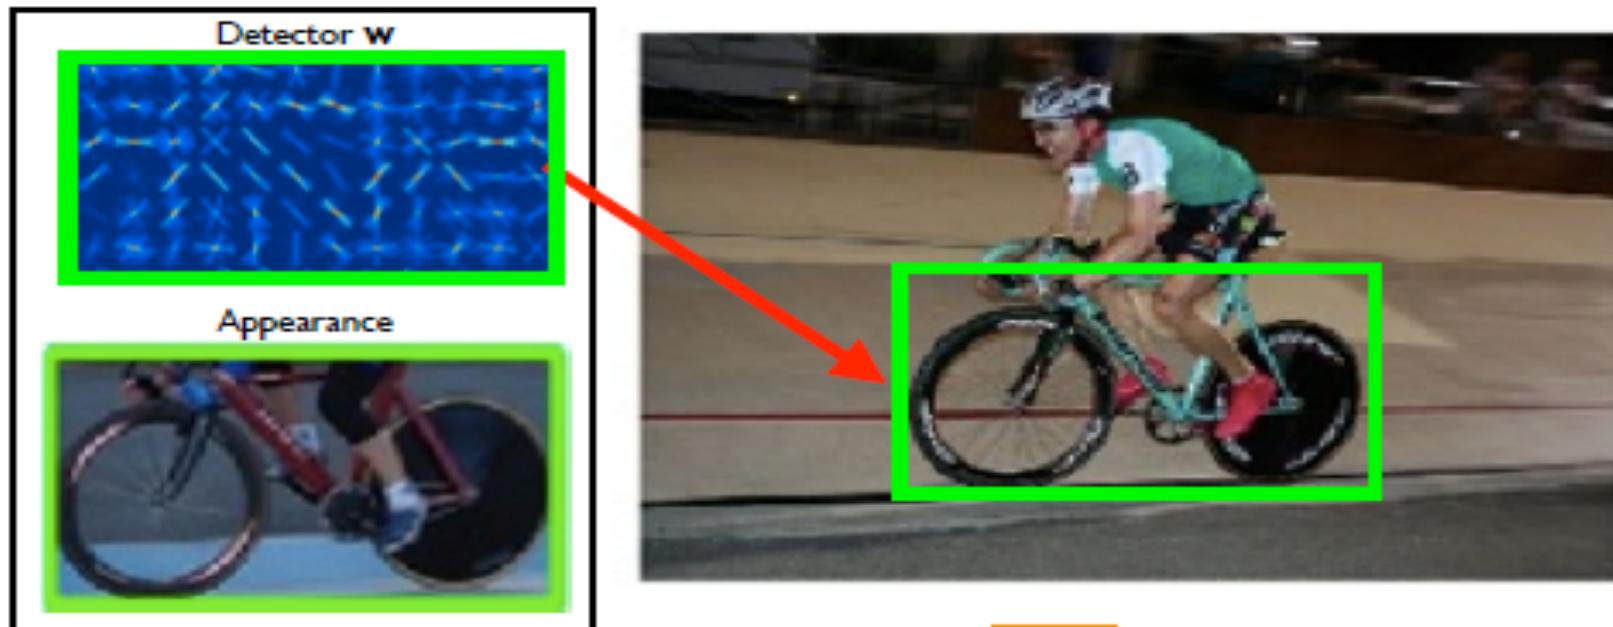
Exemplar



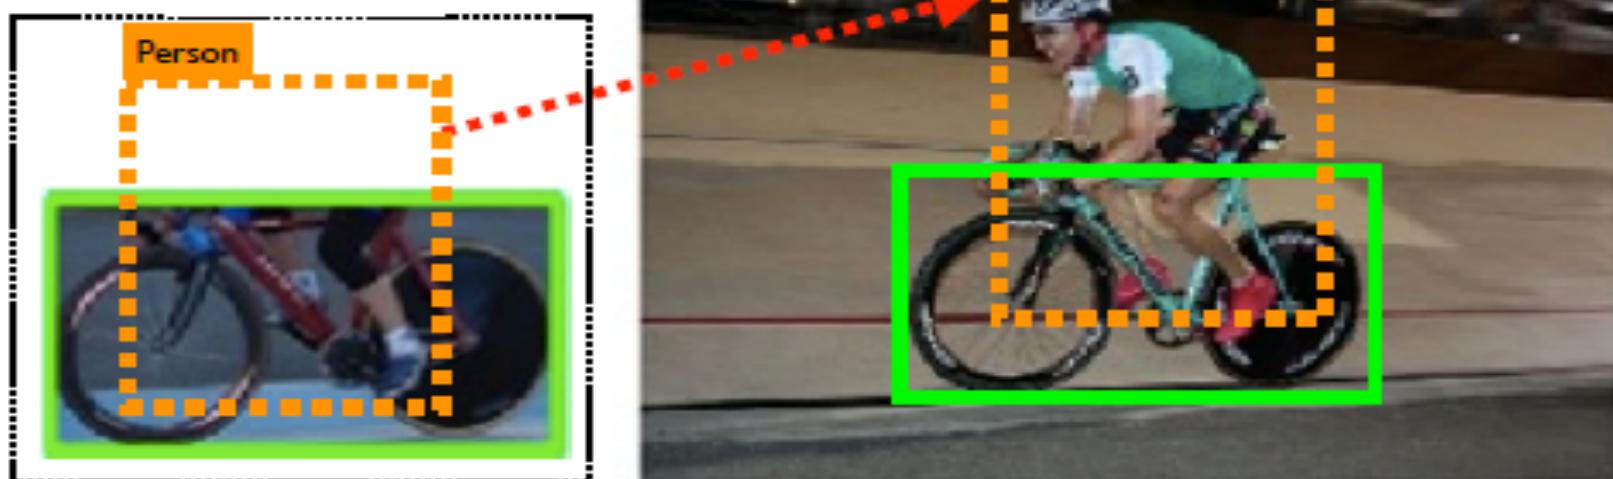
Meta-data



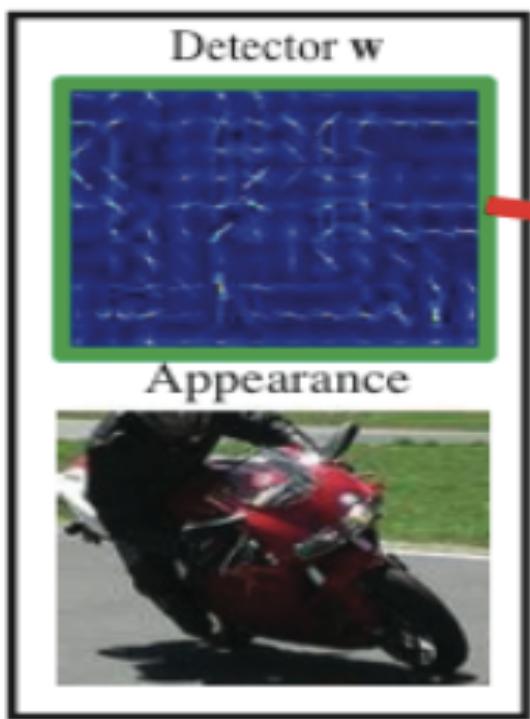
Exemplar



Meta-data



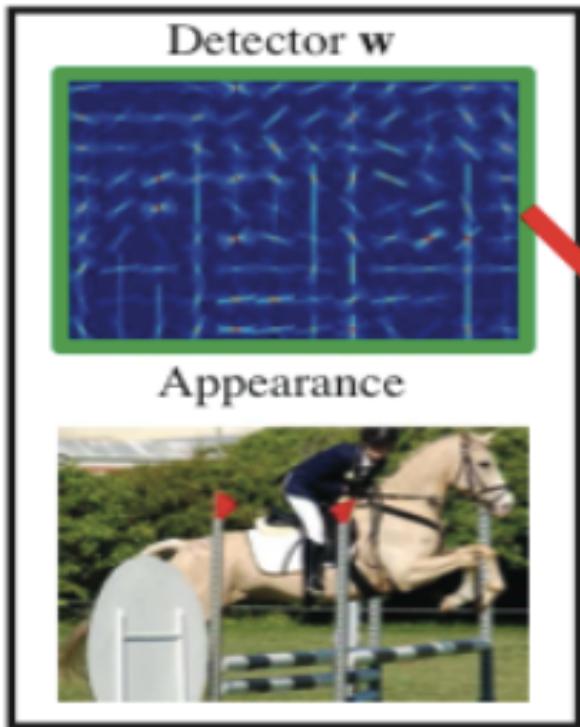
Exemplar



Meta-data



Exemplar



Meta-data



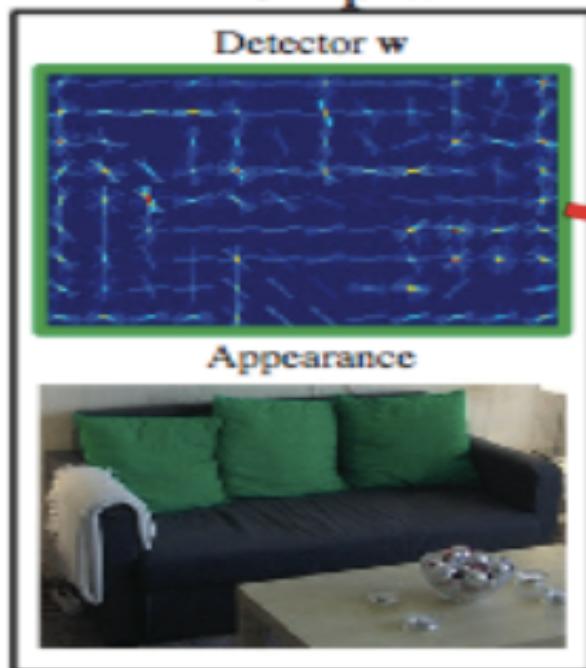
Related Object Priming: Evaluation on Person Prediction

Category	Majority Voting	us
bicycle	63.4%	72.8%
motorbike	50.0%	67.4%
horse	62.6%	77.2%

Table 2. **Is there a person riding this horse?** We predict from our bicycle, motorbike, and horse detectors whether there is a person riding the object. Our approach is better than the majority vote baseline, suggesting that exemplars are useful at predicting nearby, related objects.

More Transfer Examples

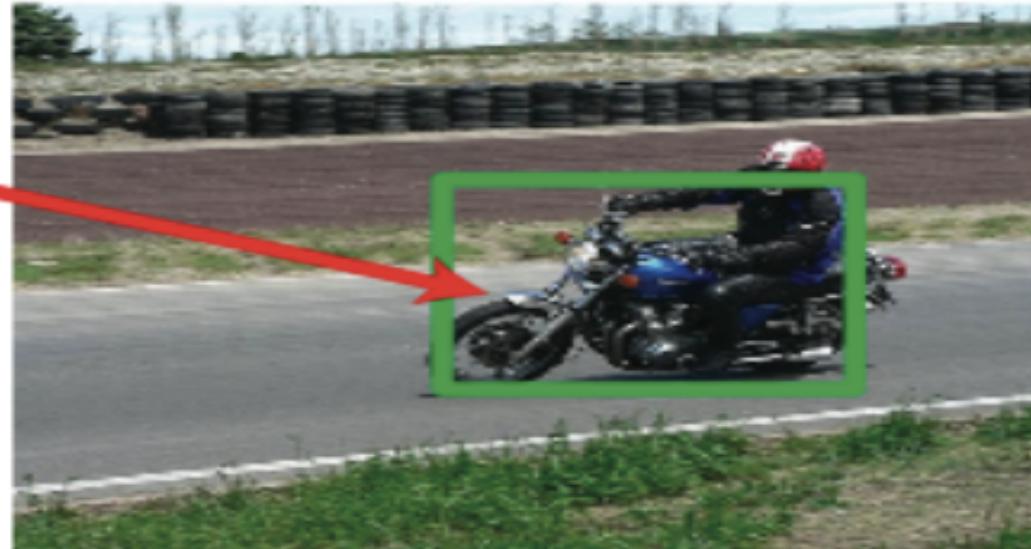
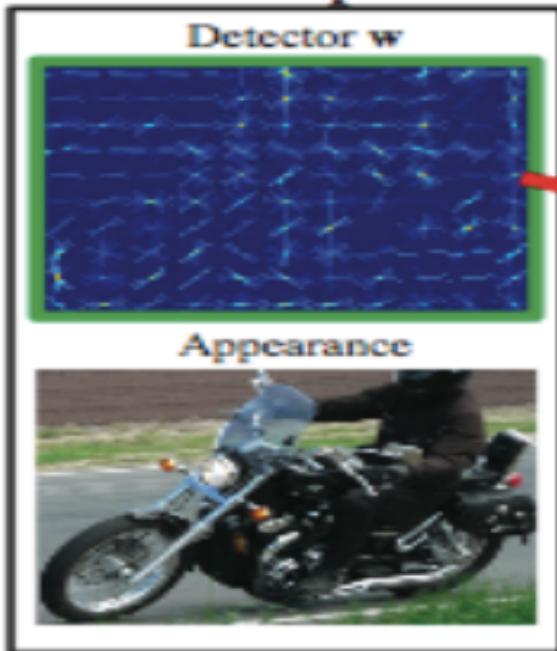
Exemplar



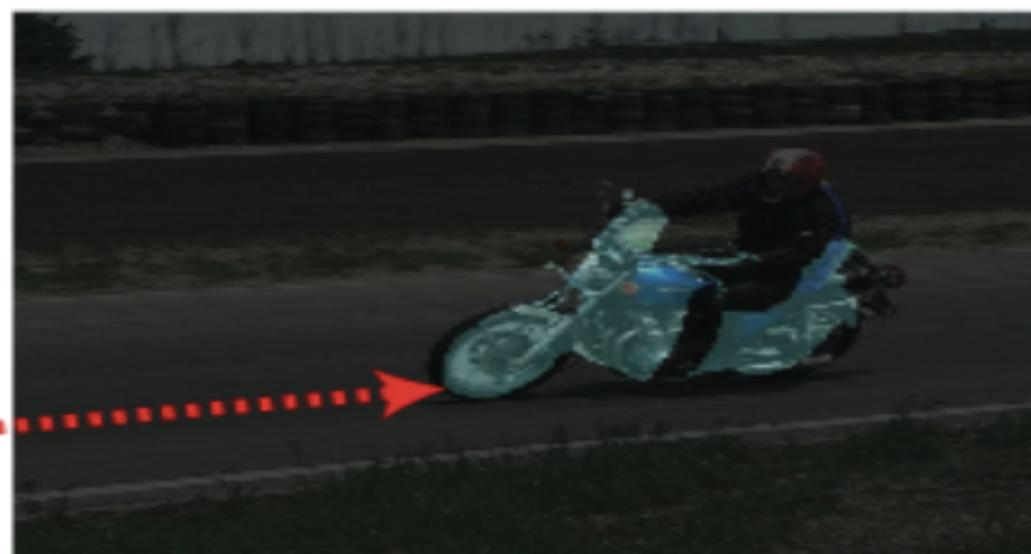
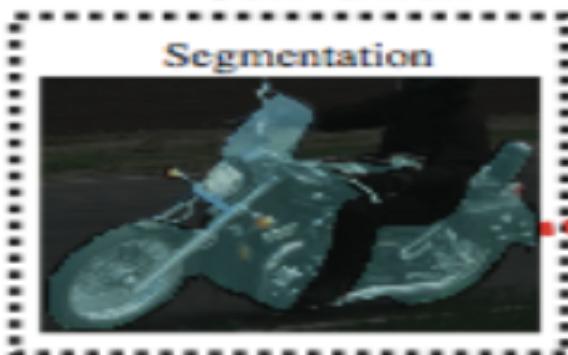
Meta-data



Exemplar

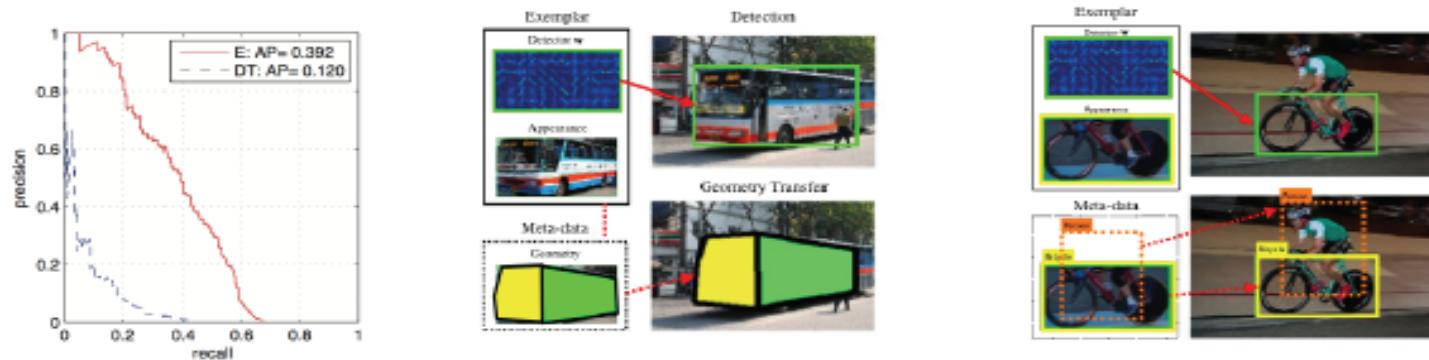


Meta-data

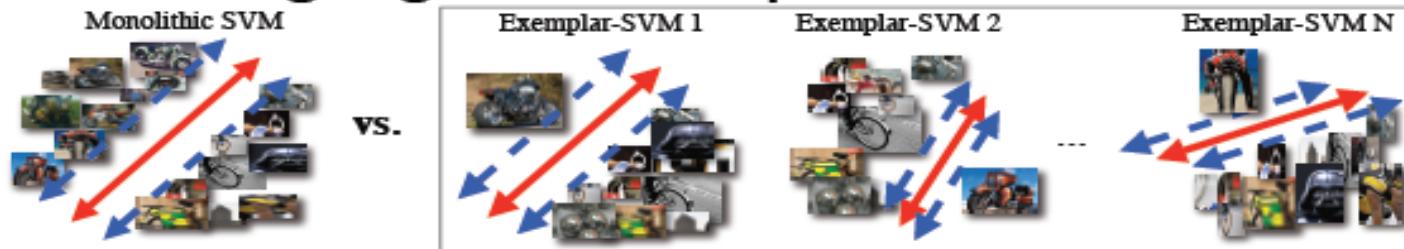


Conclusion

- ExemplarSVMs can be used for recognition, label transfer, and complementary object prediction



- Large-scale negative mining is the **key** to learning a good ExemplarSVM



- Generalization is possible from a single positive example and a vast set of negatives

!

References

- Ensemble of Exemplar-SVMs for Object Detection and Beyond, T. Malisiewicz, A. Gupta, A. Efros, ICCV 2011.
- <https://speakerd.s3.amazonaws.com/presentations/11f1d950241d013023151231381d9c14/slides.pdf>
- http://www.cs.cmu.edu/~tmalisie/projects/iccv11/malisiewicz_icml2012_talk.pdf
- <http://www.cs.cmu.edu/~tmalisie/projects/iccv11/exemplarsvm-mit-talk.pdf>